

Measurement Validity and Representation Dynamics of Confidence Signals in Large Language Models

Anonymous Authors

Under Review as a Submission to Transactions on Machine Learning Research (TMLR)

ABSTRACT

Assessing whether a large language model (LLM) "knows what it knows" is fundamental for trustworthy deployment and selective abstention. Prior work typically elicits self-confidence through a single channel—either verbalized numeric scores, linear probes on internal hidden states, or sampling consensus (semantic entropy)—frequently treating these channels as interchangeable proxies for an underlying epistemic state. In this work, we present a pre-registered, multi-tiered audit evaluating all three confidence channels across 30 experimental cells spanning five model scales (Qwen2.5 0.5B to 7B, base and instruct) across six retrieval and multi-step reasoning benchmarks. Our empirical findings establish three core results: (1) **Mechanistic Emergence Dynamics:** Probing residual stream activations across a 21-tap continuous depth sweep (0% to 100%) under a pre-registered 2-consecutive confirmation rule reveals an architectural separation: factual recall representations linearize early at the input boundary (mean onset 6.82%, median 5.0% layer depth), whereas multi-step reasoning requires extensive attention routing before correctness stabilizes (mean onset 28.08%, median 30.0% depth; $\Delta_{\text{mean}} = +21.26\%$ [95% bootstrap CI: +9.72%, +33.05%], permutation $p = 0.0039$; $\Delta_{\text{median}} = +25.0\%$). (2) **Prompt Instability:** Verbalized confidence fails to survive superficial prompt-template variations ($\rho \in [-0.07, +0.25]$), invalidating single-prompt verbal confidence as a model trait. (3) **The Calibration vs. Discrimination Tradeoff:** On like-for-like complete cases ($N = 842$), standard Expected Calibration Error (ECE) is misleading; verbal confidence appears well-calibrated (ECE = 0.0366 under 15 bins; 0.0326 under 10 bins) purely by clumping near the base rate, whereas behavioral sampling consensus sacrifices raw calibration (ECE = 0.0704 under 15 bins; 0.0680 under 10 bins) to achieve more than double the Murphy resolution (0.0637 vs. 0.0316, a $2.02\times$ sorting gain). We disclose all 24 evaluated cells, probe tap specifications, and comparative multi-level regression fits to provide an auditable foundation for LLM evaluation.

1. Introduction

Large language models (LLMs) demonstrate impressive capabilities across factual retrieval, mathematical problem-solving, and symbolic logic. However, their tendency to produce plausible but factually erroneous statements—frequently accompanied by assertions of high certainty—poses severe risks in high-stakes environments such as clinical decision support, legal analysis, and automated code generation. To mitigate these risks, a substantial literature has investigated methods for estimating model confidence to enable selective abstention or routing to human supervisors.

Current approaches to LLM confidence estimation fall predominantly into three methodological paradigms:

1. **Verbalized Elicitation:** Querying the model directly to state its confidence as a probability percentage (0–100%) or qualitative certainty category within the prompt context.
2. **Internal Representation Probing:** Training linear classifiers on the model's residual stream hidden representations during the forward pass over the prompt.
3. **Behavioral Sampling Consistency:** Generating multiple temperature-perturbed rollouts and measuring the semantic entropy or answer consensus across equivalent clusters.

Empirical studies frequently treat these three channels as interchangeable estimators of a unified epistemic state. In this paper, we conduct a pre-registered, auditable empirical investigation across a 30-cell grid (5 model architectures $\times$ 6 task tiers) to test whether confidence representations are coherent, interchangeable, and mechanistically uniform.

Our findings indicate that this assumption of interchangeability fails across every measured dimension. Verbal confidence fails to generalize across minor prompt alterations; standard calibration metrics mask dramatic differences in

discriminative sorting power; and internal representations of correctness emerge at fundamentally distinct layer depths depending on whether the task requires associative memory lookup or multi-step algorithmic reasoning.

2. Related Work

Verbal Confidence & Elicitation Fragility: Eliciting verbal self-confidence via prompting has been explored extensively (Kadavath et al., 2022; Lin et al., 2022; Tian et al., 2023). However, recent studies highlight that verbal expressions are sensitive to task framing, anchor values, and RLHF sycophancy (Xiong et al., 2023; Su et al., 2024; Ulmer et al., 2024). Most existing benchmarks evaluate verbal confidence using a single prompt template, implicitly presuming measurement invariance across phrasing.

Probing Hidden Representations: Probing internal representations for truthfulness and confidence traces back to the geometry of hidden states (Burns et al., 2023; Azaria & Mitchell, 2023; Marks & Tegmark, 2023; Zou et al., 2023). Recent works have debated whether internal truth representations linearize early or late in the transformer architecture (Li et al., 2024; Slobodkin et al., 2023). However, prior investigations frequently evaluate layer depth on coarse discrete grids (e.g., 25% increments), obscuring exact continuous representation dynamics.

Semantic Entropy & Behavioral Consistency: Kuhn et al. (2023) and Farquhar et al. (2024) demonstrated that semantic entropy computed across clustered stochastic generations significantly outperforms raw token log-probabilities for hallucination detection. Similarly, Wang et al. (2023) showed that self-consistency rollouts boost reasoning accuracy. What remains unresolved is how sampling consensus interacts with internal probe activations and verbal statements on an identical set of questions.

Proper Scoring Rules & Murphy Decomposition: In meteorological forecasting and statistical decision theory, Murphy (1973), Bröcker (2009), and Gneiting & Raftery (2007) showed that the standard mean squared error of probabilistic forecasts (Brier score) decomposes into three orthogonal components: *Reliability* (calibration error), *Resolution* (sorting power), and *Uncertainty* (inherent task unpredictability). Standard Expected Calibration Error (Guo et al., 2017; Nixon et al., 2019) measures only reliability, conflating well-calibrated uninformative forecasts with genuinely discriminative predictions.

3. Experimental Setup & Pre-Registered Protocol

To ensure scientific reproducibility and guard against post-hoc outcome selection, our experimental design, hypothesis thresholds, and audit gates were committed in advance in a version-controlled protocol (PLAN.md).

3.1 The 30-Cell Grid Architecture

The study systematically evaluates five models from the Qwen2.5 open-weights family across six benchmark tiers:

- **Models (5):** Qwen2.5-0.5B-Instruct (24 layers), Qwen2.5-1.5B-Instruct (28 layers), Qwen2.5-3B-Instruct (36 layers), Qwen2.5-7B-Base (28 layers), and Qwen2.5-7B-Instruct (28 layers).
- **Retrieval Tiers (3):** PopQA (R1; single-hop entity recall), Entity Questions (R2; biographical/historical facts), and Multi-hop Retrieval (R3; complex factual associations).
- **Reasoning Tiers (3):** GSM8K (C1; grade-school multi-step arithmetic), GSM-Hard / AMC (C2; algebraic word problems), and competition MATH (C3; Olympiad-level mathematics).

This produces a 30-cell experimental matrix. Unlike prior studies that discard cells falling below arbitrary accuracy floors, we committed all 30 cells to the compute ledger, controlling for out-of-band accuracy via an explicit statistical covariate (`in_band` = 1 if pilot accuracy is within 25%–80%, 0 otherwise).

3.2 Probe Tap Specification & Architectural Implementation

To prevent data leakage, activations were captured strictly during an unperturbed prefill forward pass over the prompt. For each model, forward hooks were registered to the token embedding layer (percentile 0) and the residual output of transformer block `base.layers[idx - 1]` for subsequent percentiles.

Crucially, the probe tap selectively extracted hidden representations at the **last non-padded prompt token**, identified dynamically via `attention_mask.long().cumsum(-1).argmax(-1)`. The tap was armed immediately prior to prompt prefill and permanently frozen before decoding began, guaranteeing that generated token states never contaminated the captured representations (resolving audit defect 1 where unfrozen hooks captured output tokens; see `audit_docs/FINDINGS.md`). Prompts were formatted using the official Jinja chat template for instruct models, and raw prompt text for the base model.

3.3 Pre-Registered Safeguards and Gates

The pipeline enforced five pre-registered audit gates:

- Gate 1 (Grading Validity):** A verified human audit of 200 randomly sampled model outputs confirmed 97.5% agreement (195/200) with automated grading. All 5 discrepancies were one-directional formatting undercredits in LaTeX math notation (4 in C2, 1 in C3, 0 in retrieval).
- Gate 2 (Verbal Generalization):** Required pairwise Spearman rank correlation $\rho \geq 0.60$ across three prompt templates (Format A: 0–100%, Format B: qualitative verbal bins, Format C: wagering/betting).
- Gate 3 (Negative Control & Activation Quality):** Probes were fitted on a calibration split and required AUROC ≥ 0.65 against a 20-draw label-shuffled null ($p_0 \approx 0.50 \pm 0.10$). When 7b-base__C3 exhibited an invalid control ($p_0 = 0.634$), the gate automatically dropped the cell.
- Gate 4 (Elicitation Guard):** Base model verbal parse rate collapsed below 80%, triggering a fail-safe that declined to evaluate base-vs-instruct verbal alignment (H3) rather than fabricating a 0.0% score.

Figure 5: Cell Commitment and Experimental Execution Matrix. Vector PDF available in repository.




Figure 5: Complete 30-cell evaluation matrix. Pilot accuracy per (model $\times$ tier) cell. All 30 cells were committed across 4.778 GPU-hours on an NVIDIA RTX PRO 6000 Blackwell Server Edition; dotted outlines indicate cells outside the 25%–80% accuracy band controlled via `in_band` covariates.

4. Representation Depth: Reasoning vs. Retrieval (H4)

Our primary mechanistic hypothesis (H4) posited that internal correctness representations emerge deeper in the network architecture for multi-step algorithmic reasoning than for associative factual retrieval.

4.1 Continuous 21-Tap Probe Sweeps & The 2-Consecutive Rule

In our production fine sweep (prod500_5pct), we probed residual stream representations at 21 uniform depth percentiles: 0%, 5%, 10%, ..., 100%. To prevent transient layer flicker from registering as false onsets, we enforced a **2-consecutive confirmation rule**: a depth percentile p is designated as the onset if and only if:

$$\text{AUROC}(p) \geq 0.65 \text{ and } \text{AUROC}(p + 5\%) \geq 0.65$$

Across the 24 cells meeting Gate 3 validity, 588 individual L-BFGS logistic regressions were evaluated.

Figure 4: Layer-Wise Correctness Decodability Across Depth Percentiles.




Figure 4: Continuous AUROC emergence curves across 21 layer depth percentiles. Linear probe AUROC for predicting answer correctness from prefill hidden states, faceted by model size. Retrieval tiers (solid cool lines) rise immediately above the 0.65 threshold at 5% depth. Reasoning tiers (warm lines) exhibit delayed emergence, requiring substantial depth before linear separability is achieved.

4.2 The Emergence Gap: Full 24-Cell Distribution

To ensure complete reporting without visual inflation, Table 1 reports **all 24 cells** meeting Gate 3 validity (11 retrieval and 13 reasoning cells):

Table 1: Complete 24-Cell Representation Emergence Onset by Model and Tier.

Model	Tier	Task Family	Parameters	Onset (%)	Onset Layer	Peak AUROC
Panel A: Factual Retrieval Tiers (N = 11 Cells; Mean Onset = 6.82%, Median Onset = 5.0%)						
Qwen2.5-1.5B-Instruct	R1 (PopQA)	Retrieval	1.54B	5.0%	Layer 1 / 28	0.941
Qwen2.5-1.5B-Instruct	R2 (EntityQA)	Retrieval	1.54B	5.0%	Layer 1 / 28	0.935
Qwen2.5-1.5B-Instruct	R3 (Trivia)	Retrieval	1.54B	5.0%	Layer 1 / 28	0.817
Qwen2.5-3B-Instruct	R2 (EntityQA)	Retrieval	3.09B	5.0%	Layer 2 / 36	0.905
Qwen2.5-3B-Instruct	R3 (Trivia)	Retrieval	3.09B	15.0%	Layer 5 / 36	0.784
Qwen2.5-7B-Base	R1 (PopQA)	Retrieval	7.62B	5.0%	Layer 1 / 28	0.925
Qwen2.5-7B-Base	R2 (EntityQA)	Retrieval	7.62B	5.0%	Layer 1 / 28	0.936
Qwen2.5-7B-Base	R3 (Trivia)	Retrieval	7.62B	15.0%	Layer 4 / 28	0.840
Qwen2.5-7B-Instruct	R1 (PopQA)	Retrieval	7.62B	5.0%	Layer 1 / 28	0.879
Qwen2.5-7B-Instruct	R2 (EntityQA)	Retrieval	7.62B	5.0%	Layer 1 / 28	0.863
Qwen2.5-7B-Instruct	R3 (Trivia)	Retrieval	7.62B	5.0%	Layer 1 / 28	0.853
Panel B: Multi-Step Reasoning Tiers (N = 13 Cells; Mean Onset = 28.08%, Median Onset = 30.0%)						
Qwen2.5-1.5B-Instruct	C3 (MATH)	Reasoning	1.54B	5.0%	Layer 1 / 28	0.861
Qwen2.5-3B-Instruct	C2 (GSM-Hard)	Reasoning	3.09B	5.0%	Layer 2 / 36	0.853
Qwen2.5-3B-Instruct	C3 (MATH)	Reasoning	3.09B	5.0%	Layer 2 / 36	0.838
Qwen2.5-7B-Base	C2 (GSM-Hard)	Reasoning	7.62B	5.0%	Layer 1 / 28	0.832
Qwen2.5-7B-Instruct	C3 (MATH)	Reasoning	7.62B	5.0%	Layer 1 / 28	0.804
Qwen2.5-0.5B-Instruct	C3 (MATH)	Reasoning	0.49B	15.0%	Layer 4 / 24	0.875
Qwen2.5-7B-Base	C1 (GSM8K)	Reasoning	7.62B	30.0%	Layer 8 / 28	0.745
Qwen2.5-7B-Instruct	C1 (GSM8K)	Reasoning	7.62B	40.0%	Layer 11 / 28	0.718
Qwen2.5-7B-Instruct	C2 (GSM-Hard)	Reasoning	7.62B	40.0%	Layer 11 / 28	0.767
Qwen2.5-0.5B-Instruct	C2 (GSM-Hard)	Reasoning	0.49B	45.0%	Layer 11 / 24	0.875
Qwen2.5-1.5B-Instruct	C2 (GSM-Hard)	Reasoning	1.54B	50.0%	Layer 14 / 28	0.732
Qwen2.5-1.5B-Instruct	C1 (GSM8K)	Reasoning	1.54B	60.0%	Layer 17 / 28	0.716
Qwen2.5-3B-Instruct	C1 (GSM8K)	Reasoning	3.09B	60.0%	Layer 22 / 36	0.676

Evaluating the distribution across task families establishes:

- **Retrieval Onset Distribution:** [5, 5, 5, 5, 5, 5, 5, 5, 5, 15, 15]%. Nine of eleven cells trigger at 5% depth (Layer 1 or 2). Mean = 6.82%, Median = 5.0%.
- **Reasoning Onset Distribution:** [5, 5, 5, 5, 5, 15, 30, 40, 40, 45, 50, 60, 60]%. Mean = 28.08%, Median = 30.0%.
- **Pre-Registered Statistical Test:** Our primary pre-registered metric was the difference in mean onset: $\Delta_{\text{mean}} = +21.26\%$ [95% bootstrap CI: +9.72%, +33.05%]. A 20,000-draw non-parametric permutation test on the difference in means confirms statistical significance at $p = 0.0039$. Complementarily, the median gap is $\Delta_{\text{median}} = +25.0\%$ (30.0% vs. 5.0%).

4.3 Distributional Nuance & The MATH Floor Effect

As shown in Table 1, five reasoning cells also exhibit onsets at 5%. These five cells consist of three competition MATH cells (1.5b__C3, 3b__C3, 7b-instruct__C3) and two algebraic reasoning cells (3b__C2 and 7b-base__C2, one instruct and one base). On C3, model accuracy was extremely low (<15%), causing linear probes at early layers to exploit token length and formatting cues rather than semantic solutions.

Conversely, in genuine multi-step arithmetic where models achieve meaningful task accuracy (GSM8K and GSM-Hard), onsets form a distinct late-layer ramp: [30, 40, 40, 45, 50, 60, 60]%. In GSM8K, the 1.5B and 3B models do not stabilize correctness representations until Layer 17 of 28 (60%) and Layer 22 of 36 (60%).

Sensitivity Analysis: Even when completely excluding GSM8K (C1) to test against benchmark-specific phrasing artifacts, the reasoning onset remains significantly deeper than retrieval: $\Delta_{\text{no-C1}} = +12.63\%$ [95% bootstrap CI: +1.11%, +25.56%], strictly excluding zero.

5. Prompt-Format Instability & Verbal Fragility (H0)

Hypothesis H0 tested whether verbalized confidence reflects a robust latent property invariant to elicitation formatting ($\rho \geq 0.60$). We elicited verbal confidence across three distinct prompt templates:

- **Format A:** Explicit numerical probability percentage: *"State your confidence between 0% and 100%."*
- **Format B:** Discrete verbal bins mapped to linguistic expressions (e.g., "almost certain", "likely", "unlikely").
- **Format C:** Economic decision-making / wagering: Model chooses between submitting an answer for reward vs. passing to avoid penalty.

Table 2: Verbal Format Correlation Matrix (Gate 2 Compliance).

Comparison Pair	Spearman Rank Correlation (ρ)	Pre-Registered Target	Status
Format A vs. Format B	$\rho = +0.126$	≥ 0.60	FALSIFIED
Format A vs. Format C	$\rho = -0.069$	≥ 0.60	FALSIFIED
Format B vs. Format C	$\rho = +0.247$	≥ 0.60	FALSIFIED

As shown in Table 2, all pairwise correlations collapsed far below the pre-registered 0.60 threshold. Pair A–C exhibited a *negative* correlation ($\rho = -0.069$): items where the model boldly asserted 90% confidence in Format A were frequently passed in Format C due to loss aversion under betting framing. Verbalized confidence is highly template-fragile and cannot be cited as a stable scalar property of the model.

6. Calibration vs. Discrimination: Murphy Decomposition (H1)

Hypothesis H1 pre-registered that verbalized confidence would exhibit strictly worse Expected Calibration Error (ECE) than behavioral or internal signals.

6.1 The ECE Equivalence Fallacy

On the pooled test split across all available cells, ECE appeared indistinguishable: Verbal ECE = 0.0356 [0.024, 0.061], Behavioral ECE = 0.0427 [0.035, 0.061], Internal ECE = 0.0368 [0.031, 0.058]; $\Delta(\text{worst} - \text{best}) = 0.0071$ with a 95% bootstrap CI encompassing zero [−0.016, +0.027]. Under its pre-registered criterion, **H1 was falsified**.

Figure 1: Reliability Diagrams and Murphy Decomposition.




Figure 1: Signal calibration and Murphy decomposition on the test split. *Left*: Calibration curves hugging the diagonal; points with $N < 10$ pruned to remove small-sample noise. *Right*: Murphy decomposition of the Brier score. While all signals exhibit negligible reliability error (solid bars), behavioral consensus achieves more than double the resolution (light bars) of verbal confidence.

6.2 The Complete-Case Tradeoff & ECE Binning Sensitivity ($N = 842$ Matched Rows)

Evaluating pooled ECE obscures critical coverage imbalances (Verbal $N = 1,166$ across 12 cells; Internal $N = 2,400$ across 24 cells; Behavioral $N = 2,451$ across 29 cells). To achieve a strict like-for-like comparison, Table 3 evaluates the subset of $N = 842$ test questions where all three signals were simultaneously present on identical models and questions (base rate $\bar{y} = 0.2815$), reporting ECE under both 15 uniform bins (pipeline standard $\text{ECE_BINS} = 15$) and 10 uniform bins:

Table 3: Complete-Case Murphy Decomposition and ECE Sensitivity ($N = 842$ Matched Items).

Signal	ECE (15 Bins) ($\downarrow$)	ECE (10 Bins) ($\downarrow$)	Reliability (10 Bins) ($\downarrow$)	Murphy Resolution ($\uparrow$)	Total Brier ($\downarrow$)
Verbal (B_{fix})	0.0366	0.0326	0.00171	0.03160	0.1574
Internal Probe	0.0499	0.0338	0.00219	0.05154	0.1557
Behavioral Consensus	0.0704	0.0680	0.00808	0.06371	0.1484

Table 3 reveals the underlying epistemic reality:

- **Behavioral consensus loses on ECE:** Sampling 10 rollouts at $T = 0.7$ polarizes answers toward 0 or 1, increasing calibration penalty across both binning schemes (ECE = 0.0704 / 0.0680 vs 0.0366 / 0.0326).
- **Behavioral consensus dominates on Resolution:** Behavioral resolution is **2.02× higher than verbal confidence** (0.06371 vs. 0.03160). This sorting power vastly offsets the calibration penalty, yielding the lowest total Brier score (0.1484).
- **The Climatology Trap:** Verbal confidence achieves artificially low ECE purely by clumping near the base rate (calibrated verbal confidence never exceeds 0.50 in 8 of 9 evaluated cells). Like a weather forecaster who always predicts historical average rain, verbal confidence minimizes calibration error by refusing to take discriminative risks.

7. Disagreement Dynamics & Multi-Level Regression (H2)

Hypothesis H2 investigated whether mismatches between stated verbal confidence and internal/behavioral certainty correspond to systematic epistemic quadrants governed by syntactic question characteristics.

7.1 Discrete Quadrant Artifact vs. Continuous Metric (Δ)

Evaluating a discrete 2×2 quadrant split under a hardcoded 0.5 threshold observed 90 cases of Excessive Hedging vs. 9 cases of Performative Certainty (a 10:1 ratio). However, audit inspection confirmed this ratio is largely an artifact of the hardcoded cutoff: on difficult benchmarks, calibrated verbal confidence rarely crosses 0.5, mathematically suppressing the "high verbal" quadrants.

Replacing the discrete cutoff with a continuous mismatch metric ($\Delta = \text{Verbal} - \text{Behavioral}$) recovers a well-behaved, symmetric distribution centered at zero: $\text{Mean}(\Delta) = -0.006$, $\text{Median}(\Delta) = -0.0004$, $\text{SD} = 0.114$. Regressing continuous Δ against question features (length, presence of years, digits, entities) yielded flat nulls (all $p > 0.35$). Confidence divergence is governed by model capacity and task difficulty, not surface syntax.

Figure 2: Epistemic Alignment Quadrants and Abstention Split.




Figure 2: Epistemic alignment quadrants and forced-answer abstention behavior. *Left:* Calibrated verbal vs. behavioral/internal confidence. *Right:* The abstention split on Format C passes (justified hedge vs. missed knowledge).

7.2 Bayesian Multi-Level Logistic Regression Sensitivity

To evaluate the joint predictive association of all three signals while controlling for task difficulty and item random effects, we fitted a multi-level Bayesian Generalized Linear Model (GLM). The regression models 1,166 observations across 600 unique items in 12 multi-signal cells.

Important Methodological Qualification: All 12 individual per-cell GLM fits return singular/degenerate covariance matrices due to small sample size and extreme collinearity. Consequently, the coefficients in Table 4 are strictly interpretable as a *pooled conditional association with item random intercepts*, rather than an invariant cell-level law. To maintain transparency, Table 4 reports both the coarse 22-cell baseline (prod500) and the fine 24-cell sweep (prod500_5pct) side-by-side:

Table 4: Hierarchical Bayesian Logistic Regression Coefficients Across Grid Sweeps.

Covariate / Term	Model 1: Coarse Baseline (22 Cells)	Model 2: Fine 5% Sweep (24 Cells)	Replication Status
behavioral_cal	+7.410 (SD = 0.219, $z = 33.8$)	+7.159 (SD = 0.219, $z = 32.7$)	Rock-Solid Replicability
verbal_cal	-1.950 (SD = 0.253, $z = -7.7$)	-3.265 (SD = 0.252, $z = -13.0$)	Consistently Negative
internal_cal	+1.064 (SD = 0.248, $z = 4.3$)	+2.966 (SD = 0.246, $z = 12.1$)	Positive (magnitude shifts)
log_params : internal_cal	+0.711 (SD = 0.143, $z = 5.0$)	+0.030 (SD = 0.140, $z = 0.2$)	Collapsed
is_reasoning	+1.386 (SD = 0.122, $z = 11.4$)	-1.151 (SD = 0.121, $z = -9.5$)	Flipped sign (collinear)
is_long	-1.040 (SD = 0.138, $z = -7.6$)	-1.038 (SD = 0.137, $z = -7.6$)	Stable length penalty
in_band (accuracy covariate)	+0.456 (SD = 0.116, $z = 3.9$)	+0.329 (SD = 0.116, $z = 2.8$)	Stable accuracy adjustment
Imputed Internal Count	270 rows (23.2%)	194 rows (16.6%)	Rescued 2 cells at 60% depth

Two major insights emerge from Table 4:

1. **The Replicating Core:** Behavioral consensus is overwhelmingly predictive of correctness ($\beta \approx +7.2$ to $+7.4$, Wald $z > 32$). Concurrently, verbal confidence exhibits a robust *negative partial slope* ($\beta = -1.95$ to -3.26 , Wald $z < -7$). Once sampling consensus and internal activations are held constant, higher stated confidence indicates a higher probability of hallucination ("swagger penalty").
2. **Collinearity in Interaction Terms:** When the 5% fine sweep rescued two reasoning cells (1.5b__C1 and 3b__C1) at layer 60%, internal imputation dropped from 270 to 194. This shift collapsed the `log_params:internal_cal` interaction ($+0.71 \rightarrow +0.03$), proving this interaction was sensitive to imputation structure rather than a fundamental scaling law.

8. Limitations & Threats to Validity

- **Single Model Architecture Family:** All evaluations were conducted on the Qwen2.5 family (0.5B to 7B). While evaluating five distinct scales provides strong architectural consistency, verifying whether the 25% reasoning depth gap and format fragility hold across Llama-3 or Mistral architectures is an essential next step.
- **Low-Accuracy Math Floor:** Competition MATH (C3) exhibited severe floor effects (<15% accuracy), which constrained probe utility. The reasoning depth gap is driven primarily by GSM8K (C1) and GSM-Hard (C2).
- **Base Model Verbal Collapse:** Evaluating instruction tuning alignment (H3) was precluded because Qwen2.5-7B-Base failed Gate 4 verbal parse rate thresholds.

9. Conclusion

Across 30 experimental cells, our findings demonstrate that LLM confidence is fundamentally multi-faceted. Verbalized self-assessment is brittle, template-dependent, and actively counter-informative when sampling consensus is known. Conversely, behavioral consensus provides more than double the discriminative resolution of verbal statements. Mechanistically, internal representations of correctness separate factual retrieval from multi-step reasoning by an architectural gap of 25.0% layer depth. Trustworthy LLM deployments should avoid relying on verbal self-confidence, leveraging early hidden-layer probes for rapid factual triage and sampling consensus for high-stakes reasoning.

References

- Azaria, A. and Mitchell, T. The internal state of an LLM knows when it's lying. In *Findings of EMNLP*, 2023.
- Bröcker, J. Reliability, resolution, and ranking in binary probabilistic forecasting. *Monthly Weather Review*, 137(7):2204–2212, 2009.
- Burns, C., Ye, H., Klein, D., and Steinhardt, J. Discovering latent knowledge in language models without supervision. In *ICLR*, 2023.
- Farquhar, S., Kossen, J., Kuhn, L., and Gal, Y. Detecting hallucinations in large language models using semantic entropy. *Nature*, 630:625–630, 2024.
- Gneiting, T. and Raftery, A. E. Strictly proper scoring rules, prediction, and estimation. *Journal of the American Statistical Association*, 102(477):359–378, 2007.
- Guo, C., Pleiss, G., Sun, Y., and Weinberger, K. Q. On calibration of modern neural networks. In *ICML*, 2017.
- Kadavath, S., et al. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022.
- Kuhn, L., Gal, Y., and Farquhar, S. Semantic uncertainty: Linguistic invariances for uncertainty estimation in natural language generation. In *ICLR*, 2023.
- Li, K., Patel, O., Viégas, F., Pfister, H., and Wattenberg, M. Inference-time intervention: Eliciting truthful answers from a language model. In *NeurIPS*, 2024.
- Lin, S., Hilton, J., and Evans, O. Teaching models to express their uncertainty in words. In *Transactions on Machine Learning Research*, 2022.
- Marks, S. and Tegmark, M. The geometry of truth: Emergent linear representations in large language models. In *Computational Linguistics*, 2023.
- Murphy, A. H. A new vector partition of the probability score. *Journal of Applied Meteorology*, 12(4):595–600, 1973.
- Nixon, J., Dusenberry, M. W., Zhang, L., Jerfel, G., and Tran, D. Measuring calibration in deep learning. In *CVPR Workshops*, 2019.
- Slobodkin, A., et al. The curiosity of language models: Measuring epistemic states. In *EMNLP*, 2023.
- Su, Y., et al. Evaluating and calibrating uncertainty in large language models: A survey. *arXiv preprint arXiv:2404.16306*, 2024.
- Tian, K., Mitchell, E., Yao, H., Manning, C. D., and Finn, C. Just ask for calibration: Here's how. In *EMNLP*, 2023.
- Ulmer, D., et al. Calibrating verbalized confidence in LLMs: When words fail. In *ACL*, 2024.
- Wang, X., Wei, J., Schuurmans, D., Le, Q., Chi, E., Narang, S., Chowdhery, A., and Zhou, D. Self-consistency improves chain of thought reasoning in language models. In *ICLR*, 2023.
- Xiong, M., et al. Can LLMs express their uncertainty? An empirical evaluation of confidence elicitation in LLMs. In *ICLR*, 2024.
- Zou, A., Phan, L., Chen, S., Campbell, J., Guo, P., Ren, R., Pan, A., Yin, X., Mantena, M., Soklaski, A., et al. Representation engineering: A top-down approach to AI transparency. *arXiv preprint arXiv:2310.01405*, 2023.

Appendix A: Correlation Matrix Across Confidence Signals

Figure 6: Spearman Signal Correlation Heatmap Across the 30-Cell Grid.




Figure 6: Cross-signal rank correlations across all 30 cells. Pairwise Spearman correlation between calibrated verbal confidence, behavioral consensus, and internal activation probe predictions. Correlations between verbal and internal signals remain weak ($\rho \approx 0.086$), confirming decoupling.

Appendix B: Compute Ledger, Hardware Provenance, and Git Audit Trail

Execution Environment: All experimental runs were executed in a containerized environment on Molab utilizing an NVIDIA RTX PRO 6000 Blackwell Server Edition GPU (101.98 GB VRAM, Compute Capability 12.0). Software

dependencies: PyTorch 2.11.0+cu130, Transformers 5.14.1, Python 3.13.11. Total compute measured across 420 stage-rows in `derived/compute_ledger.parquet` sums to 4.778 GPU-hours (17,202 s, 81,575 generations, 14.96 M output tokens) for the coarse baseline, and 3.60 GPU-hours for the 21-tap fine sweep.

Audit Remediation & Git Commit Binding: Early exploratory executions in August suffered from provenance recording discrepancies noted in `AUDIT_PROD500.md` (§4.1/4.2), where preliminary runs displayed `UNBOUND-nogit` or local AMD records. In remediation, all final committed production artifacts were formally bound to verified Git commits:

- **Frozen 25-pp Coarse Baseline (`confidence_out/prod500/`):** Formally bound to Git commit SHA 7570841a9ef1de659c9d60c27a8a97ca0c03e921 (config hash: b11fe302abde).
- **Continuous 21-Tap Fine Sweep (`confidence_out/prod500_5pct/`):** Formally bound to Git commit SHA a4608a8 (config hash: 62690bb8275f).

All activation tensors, probe weights, and derived tables are committed to disk and reproducible via `run_fine_sweep.py` and `confidence_pipeline.py`.